

Lab 3: Reducing Crime (Stage 3)

w203 Summer 2018

Madeleine Bulkow, Kim Darnell, Alla Hale, Emily Rapport

August 4, 2018

1 Introduction

As advisers to political campaigns for state office, we believe that the crime rates across the state should be of central concern to any candidate. State and local governments desire to control the crime rate in order to assure constituents that their communities are safe, and rigorous data analysis is needed to understand the factors that contribute to crime in different parts of the state. This report examines data related to crime rates provided by a current campaign in North Carolina (NC) and attempts to answer the following research question: *What variables are associated with crime rates across counties in North Carolina?*

Guided by this research question, we first examine, clean, and transform the provided data. We then propose a framework for understanding crime rates motivated by the concepts of *means*, *motive*, and *opportunity*, and link these concepts to specific variables in the data set. We then specify a model motivated by this three-part framework and examine its implications and statistical significance. Based on the findings of our primary analysis and other data exploration, we propose and interpret alternative models that provide additional predictive power. Our analysis is used to generate policy suggestions applicable to candidates seeking or defending office in North Carolina. In addition, we suggest a range of omitted variables that could form the foundation of future research on crime rates across the state.

2 Variable Definitions and Assumptions

```
# Import the data
df = read.csv("crime_v2.csv")
```

The data analyzed in this report were collected as part of a multi-year study on crime by Cornwell and Trumbull that was originally published in 1994. The data include various factors potentially related to crime for 90 of the 100 counties in North Carolina. Because of legal limitations on access to the full data set, this report will focus exclusively on the opensource data from 1987. As such, our findings and recommendations apply best to the North Carolina of the late 1980s, but may have some general application to current circumstances in the state. The list of NC counties in 1987 is identical to that for 2018, and has been consistent since roughly 1779.

The data set includes the following variables, which we present with definitions and assumptions:

county: An integer code indicating which North Carolina county a given row in the data file represents. Review of relevant factors suggests that these integers are FIPS codes, which are standard county identification codes generated by the Environmental Protection Agency (see <http://enacademic.com/dic.nsf/enwiki/49697> for details on FIPS codes for the state, including detailed maps).

year: A value of 1987 for all rows. This reflects that the current data are a subset of a larger, multi-year data set.

For the remaining values, we define them here and provide distribution summaries in **Table 1**, after we apply transformations.

crmrte: The ratio of crimes committed per person, based on reports from the Federal Bureau of Investigation's (FBI) Uniform Crime Reporting (UCR) program (see <https://ucr.fbi.gov/>).

prbarr: The ratio of arrests to offenses, based on reports from the FBI:UCR program.

prbconv: The ratio of convictions to arrests. Conviction data are taken from the prison and probation files of the North Carolina Department of Correction (NCDOC). Arrest data are based on reports from the FBI:UCR program.

prbpris: The ratio of prison sentences to convictions, taken from NCDOC.

avgsen: The average prison sentence in days; although the source is not specified, we assume these values come from the NCDOC.

polpc: The number of police officers per capita, computed using the FBI's police agency employee counts.

density: The number of people per square mile of land; we assume each unit reflects 100 people.

taxpc: The tax revenue per capita; we assume that this refers to taxes assessed in units of \$100 U.S. at the state level or lower (e.g., individual taxes, sales taxes, and property taxes).

west: An indicator code specifying whether county is in Western North Carolina (1 if yes, 0 if no).

central: An indicator code specifying whether county is in Central North Carolina (1 if yes, 0 if no).

urban: An indicator code specifying whether county is urban, defined by whether the county is in a Standard Metropolitan Statistical Area according to the U.S. Census (see <https://www.encyclopedia.com/finance/finance-and-accounting-magazines/standard-metropolitan-statistical-areas>).

pctmin80: The percentage of population that identified as belonging to a race/ethnicity other than "White" according to the 1980 U.S. Census. We note that in 1980, respondents were not permitted to select more than one racial/ethnic identification (see <http://www.pewsocialtrends.org/interactives/multiracial-timeline/> for the specific categories by census year).

mix: The ratio of face-to-face offenses (e.g., crimes targeting a person or people) to other offenses (e.g., crimes targeting property).

pctymle: The percentage of young males, defined as the proportion of the population that is male and between the ages of 15 and 24, according to the 1980 U.S. Census data.

The remaining variables in the provided data represent weekly wages in particular industries, as provided by the North Carolina Employment Security Commission (NCESC).

wcon: construction, **wtuc:** transit, utilities, and communication, **wtrd:** wholesale, retail trade, **wfir:** finance, insurance, real estate, **wser:** service industry, **wmfg:** manufacturing, **wfed:** federal employees, **wsta:** state employees, **wloc:** local government employees

The NCESC wage data do not include details regarding the number or proportion of people whose wages were included overall or by category. Moreover, it is unclear how jobs were assigned to each wage category or if there are omitted employment categories (e.g., educators, athletes, entertainers, artists). As such, we are unable to determine the representativeness or accuracy of the provided wage data.

Next, we evaluate the available data, clean it by removing anomalous values, and transform relevant variables.

2.1 Data Adjustments and Anomalies

The data set has several ratio variables, including **prbarr**, **prbpris**, **pctymle**, and **mix**, that are presented as decimal values between 0-1. To facilitate comparing the coefficients for these variables more easily with other numerical values in the data set, we converted their scale to 0-100, as in percentages.

The exception to the percentage conversion was **prbconv**, which reflects the ratio of convictions to arrests. This variable has several values that are greater than 1, which seems anomalous, if not erroneous. However, it is also the case that the **prbconv** values were calculated using numerators and denominators from different sources (i.e., the NCDOC and the the FBI:UCR program, respectively). Thus, we suspect that the extreme cases reflect a common preference among states to report criminal convictions in significant detail, even if all of the crimes for which individuals are convicted do not correspond with specific arrests that would be

reported to the federal government. Larger numbers of convictions make state and local criminal justice systems appear “tough on crime,” which tends to be popular with residents. We chose to leave the scale of this variable unmodified to avoid exaggerating any discrepancies in the measure that may have resulted from using multiple sources to calculate the values.

The variable **polpc** represents the number of police officers per known resident in a county, which is somewhat intangible on an individual scale. That is, it is awkward to refer to “0.004 police officers per person.” To address this, we multiplied the scores for this variable by 1000, facilitating descriptions such as “4 police officers per 1000 people.”

There is one county, Madison County (FIPS 115), for which the **prbarr** value is greater than 100%. This anomaly could reflect an error in data gathering or recording, but it may also reflect that it is common for individuals in this county to be arrested with greater frequency than they commit specific offenses. Given that the point of concern represents such a small part of the overall data set, could have been calculated in a logical way, and is unlikely to have an untoward impact on our analysis, we chose not to remove, replace, or adjust this score.

The data for Wilkes County (FIPS 193) are given twice in the provided data. We removed one set of these values so that they would not unduly affect the overall analysis. In addition, there are six rows in the data set that have no values for any variable. We assumed these rows were unintentionally included and removed all of them.

Data were not provided for the following counties (FIPS county codes are provided in parentheses): Camden (29), Carteret (31), Clay (43), Gates (73), Graham (75), Hyde (95), Jones (103), Mitchell (121), Tyrrell (177), and Yancey (199). We do not know why these cases were omitted from the original data set, nor can we say for certain the extent to which the omission of 1/10 counties across the state might affect the effectiveness of our analysis or recommendations. However, a review of 2012 population estimates for the omitted counties (see [HTTP://us-places.com/North-Carolina/population-by-County.htm](http://us-places.com/North-Carolina/population-by-County.htm)) indicates that all but one of them are ranked between 86-100 of the 100 North Carolina counties in overall population. The remaining omitted county is ranked 37th overall in population in the state and is close to several major metropolitan areas in the Northeast. This pattern of omission draws our attention to the variable of **density**, which we will consider with particular care.

```
# Clean the data
# Reassign the dataframe to a working variable
df_calc <- df
# Convert the prbarr, prbpris, and pctymle variables from decimals to percentages
df_calc$prbarr <- df$prbarr * 100
df_calc$prbpris <- df$prbpris * 100
df_calc$pctymle <- df$pctymle * 100
# Convert the mix variable from decimals to percentage
df_calc$mix <- df$mix * 100
# Convert the polpc variable from decimals to number of police per 1000 people
df_calc$polpc <- df$polpc * 1000
# Convert the prbconv variable from integer to numeric
suppressWarnings(df_calc$prbconv <- as.numeric(levels(df$prbconv)[df$prbconv]))
#remove row 89, which is a duplicate of row 88 (Madison County, FIPS 193)
df_clean <- df_calc[-c(89), ]
#remove rows with no data (i.e., all NA values)
df_clean <- df_clean[-c(91:97), ]

stargazer(df_clean[,3:25], title = "Summary of Descriptive Statistics for All Variables ",
          table.placement = "!h", header = FALSE)
```

Table 1: Summary of Descriptive Statistics for All Variables

Statistic	N	Mean	St. Dev.	Min	Max
crmrt	90	0.034	0.019	0.006	0.099
prbarr	90	29.524	13.767	9.277	109.091
prbconv	90	0.551	0.354	0.068	2.121
prbpris	90	41.063	8.067	15.000	60.000
avgsen	90	9.689	2.834	5.380	20.700
polpc	90	1.708	0.991	0.746	9.054
density	90	1.436	1.522	0.00002	8.828
taxpc	90	38.161	13.112	25.693	119.761
west	90	0.244	0.432	0	1
central	90	0.378	0.488	0	1
urban	90	0.089	0.286	0	1
pctmin80	90	25.713	16.985	1.284	64.348
wcon	90	285.353	47.753	193.643	436.767
wtuc	90	410.907	77.355	187.617	613.226
wtrd	90	210.921	33.870	154.209	354.676
wfir	90	321.621	53.999	170.940	509.466
wser	90	275.338	207.396	133.043	2,177.068
wmfg	90	336.033	88.231	157.410	646.850
wfed	90	442.619	59.951	326.100	597.950
wsta	90	357.740	43.294	258.330	499.590
wloc	90	312.280	28.132	239.170	388.090
mix	90	12.905	8.176	1.961	46.512
pctymle	90	8.403	2.345	6.216	24.871

3 Understanding Crime Rate

Our central goal for this analysis is to determine what variables are most clearly predictive of crime across the different counties of North Carolina. For this reason, we will use **crmrt** as our primary outcome variable.

To begin, we use Table 1 to examine the distribution of **crmrt** to determine its center and variability, based on data from 90 counties; there are no missing cases. This review reveals that the value of **crmrt** ranges from approximately 0.006 to 0.099, with a mean of approximately .034. In practical terms, this means that the 1987 crime rate in North Carolina varies from approximately 0.6 to 9.9 crimes per 100 people, with an average of 3.4 crimes per 100 people.

A histogram of the data (Figure 1) reveals that the crime rate data are positively skewed, with the majority of counties having a crime rate between 1-4% (i.e., 1-4 crimes per 100 people). The extended right tail indicates that a few counties have substantially higher crime rates, with some between 8-10% (i.e., 8-10 crimes per 100 people).

```
hist(df_clean$crmrt,
     main="Crime Rate per Person for NC Counties, 1987",
     xlab= "Crimes Committed per Person",
     ylab= "Number of Counties",
     cex.main = 1, cex.lab=.75, cex.axis = .75)
```

Although constituents are concerned with the crime rate in general, our experience suggests that they are often more concerned about crimes that result in physical harm to people (i.e., “personal crimes”, with a focus on “violent crimes”) than those that simply result in loss or damage to property (i.e., “property crimes”). For this reason, the conscientious political candidate must not simply focus on policies for reducing the general crime rate, but must consider how to perceptibly reduce the personal crime rate – and especially the violent

Crime Rate per Person for NC Counties, 1987




Figure 1: Histogram of the dependent variable, crime rate.

crime rate – for their constituents.

In the current data set, we may generally access the distinction between personal and property crime by examining the different rates of “face-to-face crime”, which reflects personal and violent crimes, and “other” crime, which includes property crime and other crime not directly involving people. This permits us to address a corollary to our primary research question, namely: *What are the variables associated with the **face-to-face** crime rates across counties in North Carolina?*

Extracting the face-to-face crime rate involves manipulations on **crmrte** involving **mix**, which the reader will recall is the ratio of face-to-face crimes to other crimes. Specifically, we begin by assuming that the total crime rate equals the face-to-face crime rate + the other crime rate. A somewhat tortuous manipulation, detailed below, allows us to calculate the ratio of face-to-face crimes among all crimes committed for each country in the data set.

$$\frac{\text{face-to-face}}{\text{total}} = 1 - \frac{\text{other}}{\text{total}} \quad (1)$$

$$= 1 - \frac{\text{other}}{\text{face-to-face} + \text{other}} \quad (2)$$

$$= 1 - \frac{1}{\frac{\text{face-to-face} + \text{other}}{\text{other}}} \quad (3)$$

$$= 1 - \frac{1}{\frac{\text{face-to-face}}{\text{other}} + 1} \quad (4)$$

$$= 1 - \frac{1}{\text{mix} + 1} \quad (5)$$

We use the resulting fraction, multiplied with the value for **crmrte** for a given county, to generate the face-to-face crime rate (i.e., **ftfcrmrte**) for that county. Below, we summarize the distribution of face-to-face crime across the counties of North Carolina in two histograms. On the left, the number of face-to-face crimes per person across counties. On the right, the proportion of face-to-face crimes to total crimes across counties.

```
# Calculate the face-to-face crime rate and examine distribution
df_clean$ftfcrmrte <- df_clean$crmrte * (1-1/(df_clean$mix+1))
```


```
summary(df_clean$ftfcrmrte)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00503 0.01886 0.02692 0.03048 0.03649 0.09343
```

```
# Plot the results
```

```
layout(matrix(c(1,2), 1, 2, byrow = TRUE))
```

```
hist(df_clean$ftfcrmrte,
     main="Face-to-face Crimes per Person\nfor NC Counties, 1987",
     xlab= "Face-to-face Crimes/Person",
     ylab= "Number of Counties",
     cex.main = .75, cex.lab=.75, cex.axis = .75)
df_clean$crmrte_ratio <- 1-1/(df_clean$mix+1)
hist(df_clean$crmrte_ratio,
     main= "Face-to-face/Total Crime Ratio\nfor NC Counties, 1987",
     xlab= "Face-to-face/Total Crimes",
     ylab= "Number of Counties",
     cex.main = .75, cex.lab=.75, cex.axis = .75)
```



We note that the distribution of the face-to-face crime rate is similar, but not identical to, the general crime rate. In particular, **f2rcrmrte** has a slightly lower mean of 3% (i.e., there are an average of 3 personal crimes per 100 people), and a range of 0.5% to 9.3% (i.e., 0.5-9.3 personal crimes per 100 people) across all counties. As these numbers suggest, the similarity in **crmrte** and **ftfcrmrte** appears to reflect that in most of North Carolina, face-to-face crime makes up 80% or more of all reported crime. In turn, this means that an analysis of crime rates across the state is, in many regards, also an analysis of face-to-face crime rates. For this reason, we choose to focus our modeling primarily on the more inclusive **crmrte** variable.

4 Models and Assumptions

To answer our research question (*What variables are associated with crime rates across counties in North Carolina?*), we use ordinary least squares (OLS) regression, also known as multiple linear regression. OLS allows us to determine what associations, if any, exist among the variables in the data set. We identify a variable to be explained – the “outcome” variable – and then perform analyses involving different combinations of “explanatory” (or “predictor”) variables from the data set to see the degree to which those variables and combinations can predict the observed outcomes. Each set of calculations, involving the outcome variable and a specific combination of explanatory variables, is referred to as a “model”. Models typically build on

each other, starting with a few explanatory variables that are determined to be particularly important, and are subsequently extended by adding more explanatory variables in order to increase the predictive value relative to the outcome variable.

For the current analysis, we model the variables contributing to crime rate across the counties of North Carolina in five stages, resulting in five models. An overview of each model is provided below.

- Model 1 includes only the variables we believe to be the main predictors of crime rate (**crmrte**), guided by a framework of *means*, *motive*, and *opportunity*: percentage of young males in the population (**pctymle**), tax per capita (**taxpc**), and population density (**density**).
- Model 2 includes the factors from Model 1 as well as several others that we believe contribute meaningfully to crime rate, including location in the state (**west**), the number of police per 1000 residents (**polpc**), the ratio of arrests to offenses (**prbarr**), the ratio of convictions to arrests (**prbconv**), and the proportion of people who do not identify racially or ethnically as White (**pctmin80**).
- Model 3 builds on Model 2 by adding more information about the location of the county (**central**), the ratio of prison sentences to convictions (**prbpris**), and the average length of prison sentence (**avgsen**).
- Model 4 builds on Model 3 and adds all other explanatory variables in the data set that are not covariant with any explanatory variables already included.
- Model 5 explores whether predictors we identify for general crime rate are comparably effective for explaining the rate of face-to-face crime across counties in North Carolina by using the secondary outcome variable **ftfcrmrte**, based on **crmrte** and **mix**.

Each of our models will be assessed to determine its consistency with the following assumptions, which are standard for classic linear regression models. The statistical quality of our findings is dependent on these assumptions being met.

- Assumption 1: Linearity in parameters, such that each fit model has slope coefficients that are linear multipliers of the associated predictor variables.
- Assumption 2: Random sampling, such that the data points are independent and identically distributed.
- Assumption 3: No perfect collinearity, such that none of the variables in the sample is a constant and there is no exact linear relationship among the predictor variables.
- Assumption 4: Zero conditional mean, such that the statistical error in the model has an expected value of 0 given any values of the predictor variables.
- Assumption 5: Homoskedasticity, such that the statistical error in the model has the same variance given any value of the predictor variables.
- Assumption 6: Normality, such that the statistical error in the population is independent of the predictor variables and is normally distributed about a mean of zero.

For each of our models, we expect a classical linear model meeting Assumption 1; this is supported by the nature of our variables and the type of analysis we employ.

We also expect all of our models to meet Assumption 2, regarding random sampling, given that our data set reflects 90 of North Carolina's 100 counties, which is very close to the overall population. As a caveat, we note that 9/10 of the omitted counties are those for which 2012 population estimates are in the lowest 15% of population totals for the state. Although it is possible that these counties could have been omitted via an appropriately executed random sampling procedure, the pattern inherent in these omissions may require additional explanation or analysis, particularly if population density emerges as a useful predictor. It is relevant to note that in Northeast cities in the early 1980s, the rate of violent crime, but not property crime, was correlated with population density (see <https://www.ncjrs.gov/App/Publications/abstract.aspx?ID=99314>).

Assumptions 3-6, which are dependent on the explanatory variables involved, will be tested independently for each model. Related to this, we note that, in some cases, it may be necessary to mathematically transform some portion of the data (e.g., convert the values for a variable to their logarithmic form) in order to facilitate

an assumption being met. For example, transformations are commonly used to linearize the relationship between variables, improve homoskedasticity, or make the data more consistent with expected practice in a given scientific discipline. Beyond cases that have specific statistical or theoretical motivations, we will use the data in its original form.

For each model, we also calculate the statistical significance of each of its coefficients. This allows us to use our models for hypothesis testing. The null hypothesis for each model coefficient, β_j , is $\beta_j = 0$, meaning that the predictor variable, x_j , has no association with the outcome variable. Statistical significance, defined as t -values with probabilities (i.e., p -values) of less than 0.05, allows us to reject the null hypothesis, and state that the given predictor has a statistically significant relationship with the outcome. As a statistically conservative best practice, we use heteroskedasticity-robust standard errors for our models even if they meet the assumption of homoskedasticity.

4.1 Model 1

We frame our initial model around the concepts of *means*, *motive*, and *opportunity*, or the situational elements that typically need to be present for an individual to be found guilty of a crime. We extend this framework to the community level, proposing that communities that have more means, motives, and opportunities for crime to occur will also be those with higher crime rates.

We conceptualize the relationship between our tripartite framework and key variables from our data set as follows:

- Means: The ability and/or resources to commit a crime and likely avoid detection. We associate this element of our framework with **pctymle**, or the percentage of young males between 15-24. These are the members of the community who are most likely to have physical abilities relevant to committing crime, such as strength, speed, and agility. Where there are higher percentages of young males, we expect higher crime rates.
- Motive: Reason(s) to commit a crime, particularly against a given target. We associate this element of the framework with **taxpc**, or the tax rate per capita. As crime often involves an involuntary redistribution of resources, we anticipate the availability of resources – and perceived inequity thereof – to encourage criminal activity. Higher tax rates are typically associated with higher incomes. Given that those who want to take others' resources must go to where those resources are, we anticipate seeing more criminal activity where tax rates are higher.
- Opportunity: Occasions when crime can occur, especially when the risk of being caught is low. We associate this element of the framework with **density**, or the concentration of people in a given area. The more people who live in a defined space, the more chances there are for conflicts to develop, inequities to be observed, and self-protective subgroups (e.g., "us" and "them") to form. Moreover, with increased population comes increased likelihood of anonymity, which reduces the likelihood of identification during criminal activity and potentially lowers the social barriers for committing crime. As such, we expect areas with higher concentrations of people to have higher crime rates.

4.1.1 Evaluating the Variables

An exploratory examination of the data suggests that this framework of means, motive, and opportunity is useful with regard to crime rates in North Carolina. All three of our selected variables are strong predictors of the general crime rate, as shown in the correlation grid and scatter plots below.

```
m1vars <- df_clean[, c(3,9,10,25)]
cormat1 <- cor(m1vars)
#cormat1
corrplot(cormat1, type = "lower", method = "number",
         title = "Correlations of Model 1 Variables",
         mar = c(0,0,1,0), cex.main = .95, tl.col = "black", tl.srt = 45, cl.cex=.6)
```


Correlations of Model 1 Variables





Figure 2: Correlation grid comparing crime rate, density, tax revenue per capita, and percent young male.

```
# scatterplots of Model 1 variables
attach(df_clean)
layout(matrix(c(1,2,3), 1, 3, byrow = TRUE))
plot(density, crmrte, main = "Crime Rate vs. \nPopulation Density\n",
     xlab = "Population Density", ylab = "Crime Rate", pch = 19)
abline(lm(crmrte~density), cex.axis=.8)
plot(taxpc, crmrte, main = "Crime Rate vs. \nTax per Capita\n",
     xlab = "Tax per Capita", ylab = "Crime Rate", pch = 19, cex.axis = .8)
abline(lm(crmrte~taxpc))
plot(pctymle, crmrte, main = "Crime Rate vs. \nPercent Young Males\n",
     xlab = "Percent Young Males", ylab = "Crime Rate", pch = 19, cex.axis=.8)
abline(lm(crmrte~pctymle))
```




Next, we examine the distributions of each of our key predictor variables.

Population density (**density**):

```
summary(df_clean$density)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00002 0.54720 0.97920 1.43600 1.56900 8.82800
```

```
hist(df_clean$density,
     main="Population Density for NC Counties, 1987",
     xlab= "Population Density per Square Mile of Land (1/100) ", ylab= "Number of Counties",
     breaks = 15, cex.main = 1, cex.lab=.75, cex.axis = .75)
```



The value of **density** ranges from a score of approximately 0.002 to 880 people per square mile of land, with a mean of 145. The distribution of county densities is right skewed, with most counties having a score of 200 or fewer people per square mile of land. We note that 9/10 counties omitted from the data set are among those with smaller populations, so it is likely that our descriptive statistics are somewhat elevated.

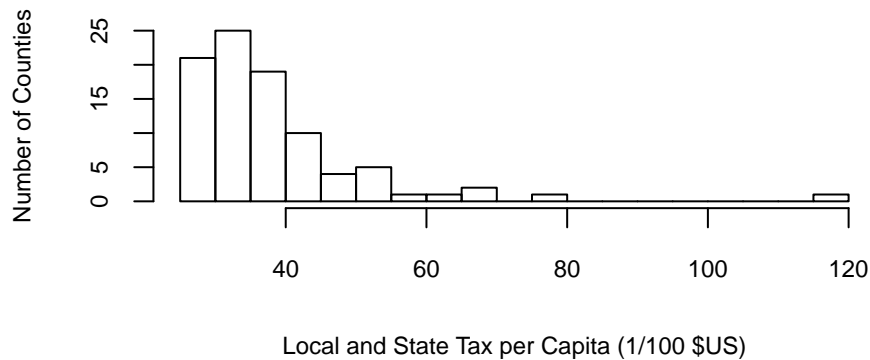
Tax per Capita (**taxpc**):

```
summary(df_clean$taxpc)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 25.69   30.73   34.92   38.16   41.01  119.80
```

```
hist(df_clean$taxpc,
     main="Tax per Capita for NC Counties, 1987",
     xlab= "Local and State Tax per Capita (1/100 $US)", ylab= "Number of Counties",
     breaks = 30, cex.main = 1, cex.lab=.75, cex.axis = .75)
```

Tax per Capita for NC Counties, 1987




The cumulative value of taxes per capita assessed at the local and state levels ranges from \$2,569 to \$11,976 per year. Once again, we see a distribution that is right skewed, with revenue in most counties below the mean of \$3,813 per year. The maximum value, the value for Dare county (FIPS 55) is nearly 50% higher than the next closest value, suggesting that this county has an anomalously high tax rate for the state. In fact, a review of the official website for Dare county (<https://www.darenc.com/>) reveals that it has an extremely active tourism industry and features a number of the state's most popular attractions, including the Outer Banks beach resort area, the Wright Brothers National Monument, the North Carolina Aquarium, and a number of other historic and recreational sites. The high rate of tax per capita for this county can be likely explained by taxes on activities related to tourism, such as those appended to hotel, rental car, and park entrance fee costs. As such, Dare county generates substantial tax-related income from people who do not live in the county, but the value of these taxes is still incorporated into the tax per capita rate for those who *do* live in the county.

Percentage of Young Males (**pctymle**):

```
summary(df_clean$pctymle)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  6.216   7.437   7.770   8.403   8.352  24.870
```

```
hist(df_clean$pctymle,
     main= "Percent Males Age 15-24 for NC Counties, 1987",
     xlab= "Percent Males Age 15-24",
     ylab= "Number of Counties",
     breaks = 30, cex.main = 1, cex.lab=.75, cex.axis = .75)
```



Across the counties of North Carolina, the percentage of males between 15-24 years of age ranges from 6.2 to 24.9. Once again, we see a distribution that is right skewed, with the majority of counties having fewer than the mean of 8.4% young males. There is one extreme value: that for Onslow county (FIPS 133). This reflects that Onslow county includes the city of Jacksonville, which contains the United States Marine Corps' Camp Lejeune and the New River Air Station, both of which are inhabited predominately by males between 18-25 years of age.

4.1.2 Evaluating Assumptions

Before we build our model, we check the correlations between our independent and predictor variables to assess the variables' consistency with Assumption 3 regarding no perfect collinearity. We can reference the correlation matrix above to see that all the correlations among our variables are less than 1. Specifically, **crmrte** is positively predicted, in descending order, by **density** ($r = .73$), **taxpc** ($r = .45$), and **pctymle** ($r = .29$). Additionally, **density** shows a strong positive correlation with **taxpc** ($r = .32$) and a weaker positive correlation with **pctymle** ($r = .12$). None of the variables are perfectly collinear with any of the others, so Assumption 3 is validated.

Given the skewed nature of **crmrte**, we explore a log transform of the dependent variable and build two versions of Model 1: one with untransformed **crmrte** and one with a log transform of **crmrte**. A comparison of the two versions (**Table 2**) and their consistency with the statistical assumptions allows us to select the better foundational model.

```
# Build Model 1 (untransformed) and Model 1 (log)
model_1 = lm(crmrte ~ density + taxpc + pctymle, data = df_clean)
model_1_log = lm(log(crmrte) ~ density + taxpc + pctymle, data = df_clean)
```

Table 2 summarizes the results of the linear regressions on the untransformed and log transformed versions of Model 1. Standard errors are given in parentheses.

```
stargazer(model_1, model_1_log, type = "latex",
  title = "Linear Regression Results for Two Versions of Model 1",
  keep.stat = c("rsq", "n", "adj.rsq"),
  se = list(sqrt(diag(vcovHC(model_1))), sqrt(diag(vcovHC(model_1_log)))),
  single.row = TRUE, table.placement = "!", header = FALSE,
  star.cutoffs = c(0.05, 0.01, 0.001))
```

The version of the model with an untransformed **crmrte** has an the adjusted R^2 value which indicates that 62.8% of the variance in crime rate can be explained by our combination of three variables. In contrast, in


Table 2: Linear Regression Results for Two Versions of Model 1

	<i>Dependent variable:</i>	
	crm rte	log(crm rte)
	(1)	(2)
density	0.008*** (0.001)	0.195*** (0.039)
taxpc	0.0004 (0.0003)	0.009 (0.007)
pctymle	0.002* (0.001)	0.055** (0.021)
Constant	-0.009 (0.013)	-4.614*** (0.328)
Observations	90	90
R ²	0.640	0.481
Adjusted R ²	0.628	0.463
<i>Note:</i> *p<0.05; **p<0.01; ***p<0.001		

the version of the model with the log transformed **crm rte**, the adjusted R^2 has dropped to 46.3%, reflecting a substantial reduction in predicted variance in the outcome variable. As such, the untransformed version of the model has greater predictive power with regard to outcome variable of interest.

Next, we evaluate the Cook's Distance for the residuals by examining the residuals vs. leverage plots for these two candidate versions of Model 1.

```
layout(matrix(c(1,2), 1, 2, byrow = TRUE))
plot(model_1, which= 5,
      caption = "", main = "Residuals vs Leverage\n for Model 1 (untransformed)",
      cex.main = .9, cex.lab=.75, cex.axis = .75)
plot(model_1_log, which = 5,
      caption = "", main = "Residuals vs Leverage \n for Model 1 (log)",
      cex.main = .9, cex.lab=.75, cex.axis = .75)
```




In both versions of Model 1, we find one point (i.e., the one corresponding to Dare county) that has a Cook's Distance greater than 1. As noted previously, Dare county has substantially higher tax per capita than other North Carolina counties, which we attribute to tax revenue from tourism. As such, the deviation of this single case is understandable and does not warrant its removal.

Assumption 4, the assumption that the mean error conditional on the dependent variable equals zero, is not met in a strict causal sense for either version of Model 1. That is, we cannot observe “true relationships” among the underlying quantities, due to the likely presence of other factors that may influence the outcome variable. For this reason, we can only validate a weaker, but still sufficient, form of the assumption: exogeneity of the error. This means that our interpretations of the relationships among the variables in our data will be strictly associative. From this point forward, Assumption 4 and Assumption 6 will be discussed in this more limited associative manner.

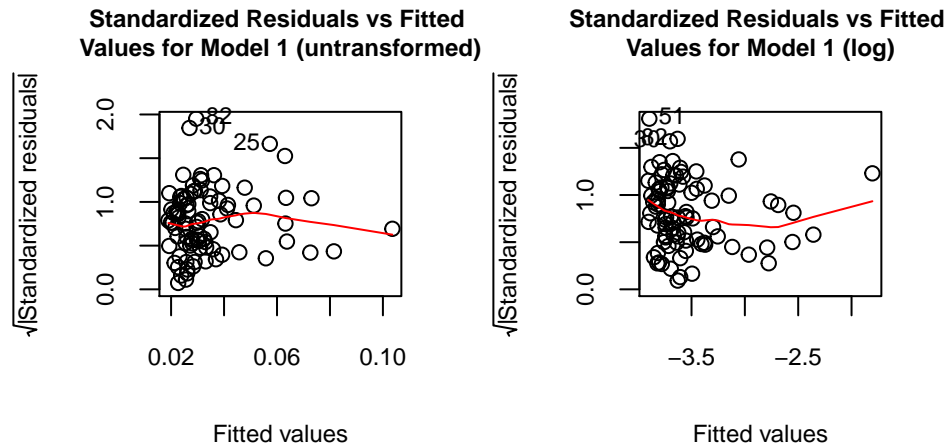
The residuals vs fitted values plots for both versions of Model 1 show that the untransformed version exhibits close to a zero-conditional mean for the errors, whereas the log transformed version has a marked curvature in the errors. This outcome preferences the untransformed version of the model.

```
layout(matrix(c(1,2), 1, 2, byrow = TRUE))
plot(model_1, which= 1,
      caption = "", main = "Residuals vs Leverage\n for Model 1 (untransformed)",
      cex.main = .9, cex.lab=.75, cex.axis = .75)
plot(model_1_log, which = 1,
      caption = "", main = "Residuals vs Leverage\n for Model 1 (log)",
      cex.main = .9, cex.lab=.75, cex.axis = .75)
```



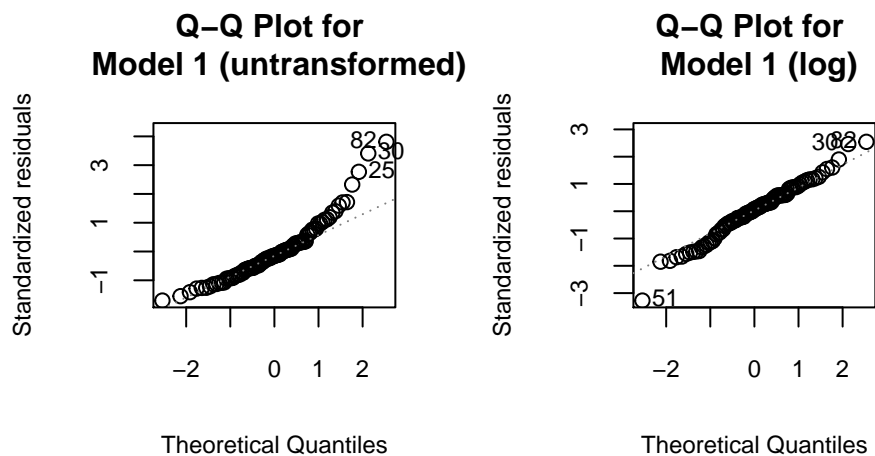
To validate the model in terms of Assumption 5 regarding homoskedasticity, we plot the standardized residuals vs. fitted values for each version. We find that the range of errors is relatively more constant throughout the range of fitted values for the untransformed model than the transformed model. We note that there are fewer data points at the higher values of **crmrte** than the lower values, so validity in this case may be somewhat weaker than with other assumptions. Nonetheless, this outcome also preferences the untransformed version of the model.

```
layout(matrix(c(1,2), 1, 2, byrow = TRUE))
plot(model_1, which= 3, caption = "", main = "Standardized Residuals vs Fitted \nValues for Model 1 (untransformed)",
      cex.main = .75, cex.lab=.75, cex.axis = .75)
plot(model_1_log, which = 3, caption = "", main = "Standardized Residuals vs Fitted \nValues for Model 1 (log)",
      cex.main = .75, cex.lab=.75, cex.axis = .75)
```



To validate Assumption 6, the normality of the residuals, we examine a Q-Q plot of the standardized residuals for each version of the model. We look for the models' standardized residuals to fall along the fit line, which is evidence that they are normally distributed, indicating that variation in the outcome variable not predicted by the model is likely due to random noise. We observe a fairly straight line with the exceptions of the tails in both versions, with the untransformed model showing substantially more deviation from the fit line at higher quantiles. Log transformation is specifically intended to normalize the distribution of residuals, so as expected, this evaluation preferences the log transformed version of the model.

```
layout(matrix(c(1,2), 1, 2, byrow = TRUE))
plot(model_1, which= 2, caption = "", main = "Q-Q Plot for \n Model 1 (untransformed)",
     cex.main = 1, cex.lab=.75, cex.axis = .75)
plot(model_1_log, which = 2, caption = "", main = "Q-Q Plot for \n Model 1 (log)",
     cex.main = 1, cex.lab=.75, cex.axis = .75)
```



Overall, our comparison of the untransformed and transformed versions of Model 1 suggests that the untransformed version is superior with regard to meeting assumptions of the classic linear model and demonstrating predictive power. That is, although the log transformed version better addresses the demands of Assumption 6, it also appears to violate Assumptions 4 and 5.

From this point forward, Model 1 will refer to the untransformed version of the model only.

4.1.3 Statistical Inference for Model 1

The matrix below summarizes the coefficients, heteroskedasticity-robust standard errors, t -values, and p -values for the parameters in Model 1. The coefficient on **density** is statistically significant at the .0001 level; the one for **pctymle** is statistically significant at a .05 level. As such, we can reject the null hypotheses that 1) there is no association between **density** and **crmrte** and 2) there is no association between **pctymle** and **crmrte**. We fail to reject the null hypothesis that there is no association between **taxpc** and **crmrte**.

```
coeftest(model_1, vcov = vcovHC)

##
## t test of coefficients:
##
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.00910806  0.01266244 -0.7193    0.4739
## density      0.00759718  0.00124324  6.1108 2.805e-08 ***
## taxpc        0.00039646  0.00029917  1.3252    0.1886
## pctymle      0.00197340  0.00082687  2.3866    0.0192 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

4.1.4 Interpreting Model 1

Our selected version of Model 1 is a level-level model, indicating neither the outcome nor any of the predictors have been transformed. For this reason, we interpret the coefficient for each predictor as describing the magnitude of the associated change in **crmrte** given a one-unit increase in the predictor, holding all other predictors constant.

As shown in Table 2, the coefficients in Model 1 are all positive. This indicates that as the population density, tax per capita, or the percentage of young males increase, an associated increase occurs in the crime rate. For example, our model predicts that, when the other predictors are held constant, a one-unit increase in **density** is associated with approximately a 0.0076 increase in the **crmrte**, and this change is statistically significant. In other words, for every additional 100 people per square mile of land, the crime rate increases by .0076; moreover, this change is too large to be attributable purely to chance. Similarly, a one-unit (1%) increase in **pctymle** is associated with a statistically significant 0.002 increase in **crmrte**. In contrast, the one-unit (\$100US) increase in **taxpc** is associated with a non-significant 0.0004 increase in **crmrte**.

Our analysis makes it clear that our three key predictor variables are not equally useful when it comes to explaining the crime rate. Specifically, increases in population density result in an increase of the crime rate that is an order of magnitude greater than that for increases in tax per capita, and about four times greater than that generated by a higher percentage of young males. The statistical significance of each coefficient maps neatly to its size: **density**'s coefficient is the highest in magnitude and has the highest statistical significance, and **taxpc**'s coefficient is the lowest in magnitude and has the lowest statistical significance. This pattern suggests that the candidate's political energy would best be spent on addressing crime-related concerns connected to population density first, followed by those associated with the proportion of young males. The tax revenue per capita, given its low explanatory power with regard to crime rate, should be of limited interest to the candidate in and of itself.

If we consider these findings in the context of crime reduction policy, a number of actionable strategies emerge. In addition, we can identify specific factors that were not present in the model or the data set, but may be useful to think about during the decision making process; these omitted variables are also **bolded** in the discussion below.

To begin, we see the **density** of an area as a contributor to the *opportunity* individuals have to commit crime. There are a number of reasons why increases in population density could facilitate increases the crime rate. For example, the more people live in a particular space, the more opportunities there are for them to come into conflict with one another, to interact with others who have different access to desirable resources, and to

be unfamiliar with or disconnected from the others with whom they come in contact day after day. Thus, to the extent that candidates have constituencies in high population areas, they should consider addressing the crime rate by developing policies that improve the ease with which large numbers of people can live and move in the same space, while reducing opportunities for conflict and perceptions of inequity. For example, **infrastructure** projects that increase the livability and communal nature of high population areas, such as well-maintained public parks and recreational areas, effective public transportation, and improved traffic and parking management, may make it easier for residents live in close quarters with others and reduce the number of negative experiences that might lead to criminal behavior. Similarly, addressing problems related to **socioeconomic inequity**, such as access to adequate educational and employment opportunities, social support programs, and affordable housing should also result in a reduction in crime, as they reduce the demand to take resources from others to fulfill basic needs. Last but not least, there is the issue of anonymity. Certainly it is easier to commit a crime against a stranger than it is a neighbor or a friend, if only because there are fewer personal costs and a lower likelihood of being recognized and/or caught. So, investing in events, facilities, and services that encourage people to get to know and develop positive relationships with those around them, take pride in their joint **community membership**, and recognize their common interests and values should also reduce crime. These might include celebrations around a famous local figure or event, neighborhood service opportunities (e.g., creating and maintaining a community garden or playground), or fundraising activities for an important local or regional cause (e.g., recovering from a natural disaster or terrorist attack).

We associate the **pctymle** in an area with the proportion of the population that has the physical *means* to commit crime. There is certainly strong cultural value in the U.S. that links masculinity with wealth and the ability to provide for a family, as well as a sense that it is acceptable – if not expected – for men to be aggressive or violent when their needs and wants are not immediately met. In fact, beliefs relevant to these values may be even more common in communities that prioritize traditional social roles, which is likely to be true in a historically conservative state such as North Carolina. To the extent that a candidate’s constituency includes communities with large percentages of young men, it could be fruitful to consider the role that local **culture** plays in young men viewing criminal behavior as a reasonable solution to their problems, and how providing constructive alternatives for young men to demonstrate and reinforce their social value could reduce crime. Relevant policies could support educational, vocational, and athletic programs, as well as involve young men in activities that contribute positively to the community and encourage them to develop it rather than damage it. Religious institutions could play an important role in this effort, as it is common for religious doctrine across faiths to discourage violence against others and other types of criminal behavior.

Our framework associates **taxpc** in an area with potential *motive* for committing crime, particularly in socioeconomically diverse communities where some “have” – sometimes to a conspicuous degree – and others “have not.” Specifically, we reasoned that areas where residents make more money would pay higher taxes *and* be more tempting targets for crime, because their higher income affords them more access to desirable items and services. However, this variable did not emerge as a statistically significant predictor in the linear regression, suggesting that we need to reconsider how the income of people in an area might be related to the likelihood that they will experience crime. In particular, if we had observed that higher income areas were also higher crime areas, we might have suggested that increased police presence and other improved security measures in those areas might dissuade people from committing crime there. Rather, we observe that the income of an area, as reflected in the tax base, does not predict crime effectively; people in low, middle, and high income areas seem to be experiencing crime at roughly similar rates. The question this raises in our mind is: Are these different communities experiencing the same kinds of crime? For example, might higher income areas experience more property crime, whereas lower income areas experience more personal crime? If communities on different rungs of the wealth ladder are experiencing reported crime to a similar degree, then policy decisions relevant to addressing the crime rate must offer solutions that serve all of these communities equitably. What will clearly *not* be helpful are policies that target members of one community as criminals and members of another as victims, or encourage the unjustified harassment, mistreatment, and arrest of members of marginalized groups in the name of reducing crime for the majority group. In fact, discriminatory policing policies are more likely to foment distrust among different communities and might even increase the damage from criminal activity by discouraging people in the targeted communities from reporting crimes. As such, **fairness of policing policy** could be an interesting variable to explore with regard to its relationship

to crime rate.

Another variable of interest suggested by our analysis is the level of **tourism**. In areas where there is a lot of tourism, there may be special circumstances that contribute to the crime rate. If tourism is seen by a community as an opportunity for good employment and additional funding for community infrastructure and beneficial social programs, we would expect this to lower the crime rate, particularly against visitors. If, on the other hand, it is viewed as an influx of distracted strangers with an abundance of extra money and little familiarity with the local environment, the crime rate would likely go up. Tourism contributes positively to the economy of North Carolina and reduces the direct tax burden on residents, while encouraging improvements in infrastructure and community facilities. To the extent a candidate can frame tourism and the support thereof as a strong positive for the their constituency (e.g., as one means to reduce local socioeconomic disparity), constituents may be less likely to want to engage in criminal activity that negatively influences the tourism industry.

4.2 Model 2

For Model 2, we include additional variables that seem to fit into our *means, motive, and opportunity* framework – albeit less neatly – and have substantial correlations with the crime rate. Of particular interest in this model is identifying other elements of the data set that relate to the *motive* dimension of our conceptualization of crime rate, as **taxpc** did not show a significant predictive effect.

We begin by hypothesizing that the more likely one is to get caught and suffer negative consequences for engaging in criminal activity, the less likely one is to want to do it. As such, we expect variables like the ratio of arrests to offenses (**prbarr**), the ratio of convictions to arrest (**prbconv**), and the number of police per capita (**polpc**) to have a negative impact on the crime rate.

We also hypothesize that socioeconomic disparity could encourage the involuntary redistribution of resources between those who have more and those who have less. To explore this, as well as the role that race/ethnicity might play in this phenomenon, we add **pctmin80** and **west** to the model. For these variables, we anticipate a positive relationship with **crmrte**, because of the socioeconomic disparity that is often associated with the communities these variables represent. That is, in areas with higher percentages of non-White minorities, conflict may arise between Whites and others if Whites are perceived by the others to be undeservedly or disproportionately privileged or Whites perceive that the others are attempting to infringe on their privilege. Similarly, in the western part of the state, there are numerous communities of lower socioeconomic status who are visited by higher status tourists from outside the region, eager to explore Appalachia and its unique culture. In this context, we can imagine that the marginalized communities of rural Appalachia might have feelings of resentment toward outsiders who view them as part of a cultural tourist attraction, whereas members of advantaged communities from outside Appalachia might feel entitled to mistreat or exploit natives to the region because they are poor. The race/ethnicity dimension emerges when we consider that minority communities are, by definition, not White, whereas the population of western North Carolina is predominantly White. Should **pctmin80** pattern differently than **west** in the model, it could suggest that marginalized communities experience crime differently, depending on their racial/ethnic composition.

Below, we display the relationships among the continuous explanatory variables. The variable **west**, which is coded as an indicator variable (i.e., each county is either “in the western part of the state” or “not in the western part of the state”), is not included. We note that for the data set as a whole, 24.4 % of all counties were coded as being in the western part of the state.

```
vars <- c("polpc", "prbarr", "prbconv", "pctmin80")
suppressWarnings(scatterplotMatrix(df_clean[,vars], diagonal = (method = "histogram")))
```

The matrix plot shows little collinearity among the considered variables, validating Assumption 3 for Model 2.

Next, we build the model and evaluate its consistency with the remaining statistical assumptions. To conserve space and avoid repetition, we will not detail all aspects of the results here, but summarize the findings for

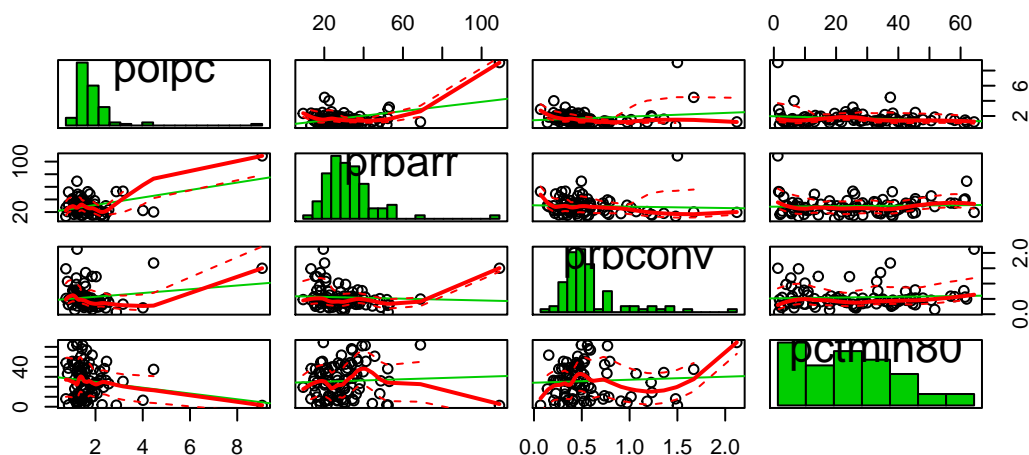


Figure 3: Scatterplot matrix showing relationships between **polpc**, **prbarr**, **prbconv**, and **pctmin80**.

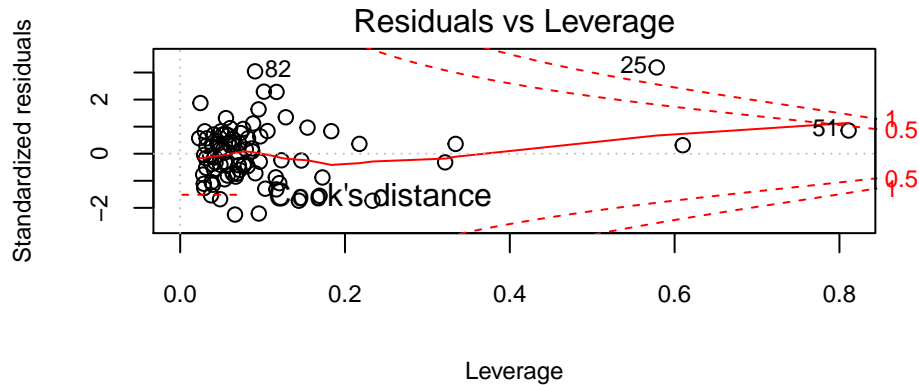
all the models in **Table 3** in section 4.6.1. here we will simply focus on the most important findings related to validating and interpreting the model.

```
# Build Model 2
model_2 = lm(crmrte ~ density + taxpc + pctymle + west + polpc + prbarr + prbconv + pctmin80,
             data = df_clean)
```

We observe that Model 2 has positive coefficients for **density**, **taxpc**, **pctymle**, **polpc**, and **pctmin80**, indicating that crime rate increases as these variables increase. Conversely, the coefficients for **west**, **prbarr**, and **prbconv** are negative, indicating that crime rate decreases as these increase. We also see that the adjusted R^2 has increased from approximately 63% to 82% between Model 1 and Model 2, indicating improved predictive power relative to **crmrate** with the addition of the new *motive*-related variables.

Now to examine the model with regard to the remaining assumptions. First, we look to see any unexpected data points display unusual influence in the new model.

```
# Influence plot
plot(model_2, which = 5,
     cex.main = 1, cex.lab = .75, cex.axis = .75)
```




(crrmrte ~ density + taxpc + pctymle + west + polpc + prbarr + prbcor

We note that point 25 still has high leverage, but as it is a single point and its status in this regard has not changed between Models 1 and 2, we are generally unconcerned about it.

Assumption 4, exogeneity, is verified below, by checking the plot of residuals against fitted values. We see a fairly flat mean across the entire range of fitted values.


```
plot(model_2, which = 1,
      cex.main = 1, cex.lab=.75, cex.axis = .75)
```



(crrmrte ~ density + taxpc + pctymle + west + polpc + prbarr + prbcor

Assumption 5 is validated for this model in the same manner as for Model 1, by examining the plot of standardized residuals vs. fitted values. As before, we find that the range of errors is relatively constant throughout the range of fitted values, with fewer data points at the higher end the fitted values' scale.

```
plot(model_2, which= 2,
      cex.main = 1, cex.lab=.75, cex.axis = .75)
```



(*crmrte* ~ *density* + *taxpc* + *pctymle* + *west* + *polpc* + *prbarr* + *prbconv*)

For Assumption 6, we note that the Q-Q plot of the standardized residuals appears much closer to linear than for Model 1, indicating that the current model includes important sources of variation that were missing from our foundational model.

Statistical Inference for Model 2

```
coefTest(model_2, vcov = vcovHC)
```

```
##
## t test of coefficients:
##
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.8840e-02  8.1882e-03  2.3009 0.0239671 *
## density      5.4313e-03  1.3879e-03  3.9134 0.0001886 ***
## taxpc        1.7816e-04  2.3922e-04  0.7447 0.4585871
## pctymle      8.7469e-04  3.9818e-04  2.1967 0.0309008 *
## west        -1.6253e-03  2.5319e-03 -0.6419 0.5227198
## polpc        6.5622e-03  2.0968e-03  3.1296 0.0024336 **
## prbarr       -5.4952e-04  1.2699e-04 -4.3272 4.279e-05 ***
## prbconv      -1.8753e-02  3.9228e-03 -4.7806 7.702e-06 ***
## pctmin80     3.2930e-04  7.8636e-05  4.1876 7.123e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The heteroskedasticity-robust *t*-test for Model 2 shows that, with the exception of **west**, all of the new variables we added to the model are statistically significant to at least a .01 level. As such, we can reject the null hypotheses that police per capita (**polpc**), arrests per offenses (**prbarr**), convictions per arrests (**prbconv**), and percentage of non-Whites (**pctmin80**) have no association with **crmrte**. Conversely, we fail to reject the null hypothesis with regard to the **west** region, such that whether a county is in the western part of the state has no predictive value with regard to the crime rate. As in Model 1, we again fail to reject the null hypothesis regarding tax per capita (**taxpc**).

Interpreting Model 2

The pattern of significance for the three foundational variables included in Model 1 has not changed. As such, our interpretation of Model 2 will focus the variables that were added at this stage.

As expected, the ratio of arrests to offenses and of convictions to arrests are negatively predictive of crime rate. This suggests that as the risk of punishment for committing crime increases, the crime rate decreases. In addition, where crime rates are lower, it is easier for law enforcement officers to more thoroughly investigate crimes, leading to more arrests and convictions. Also as expected, the percentage of non-Whites is positively predictive of crime rate, such that counties that have more racial/ethnic minorities experience higher crime rates.

Contrary to expectation, the number of police officers per capita (**polpc**) is *positively* predictive of crime rate. That is, having more police officers in an area doesn't discourage crime. Rather, the more police there are in a county, the higher the crime rate is. We interpret this to mean that the presence of more police results in more criminals being arrested and convicted for crimes that are committed, whereas fewer criminals are actually observed, arrested, or convicted in places where there are fewer law enforcement officers.

There is no meaningful statistical relationships between whether or not a county is in the western part of the state (**west**) and crime rate. This is interesting, because those of lower socioeconomic status in Appalachia are predominantly White. This contrasts with those of traditionally lower socioeconomic status in urban settings, who are typically not White. As such, the difference in significance between the predictive power of the percentage of non-White minorities and the location in the western part of the state suggests that race/ethnicity, not socioeconomic class, is more relevant to crime rate. What is not clear is if racial/ethnic minorities are more likely to be arrested and convicted of crime than their White counterparts of comparable socioeconomic status, or if there is truly more crime in minority communities than White communities, even in cases where socioeconomic status is generally held constant. Another possibility is that the popularity of tourism in the western part of the state mitigates the impact of lower socioeconomic status on crime, resulting in poorer White communities experiencing less crime in North Carolina because of where they are in the state, their contribution to White Americana, and the likelihood that most tourists are White. Given the long and complicated history of race/ethnicity and socioeconomic status in the South, we think this is an important area for future research that has powerful policy implications. We cannot, however, address all of the relevant issues here with the current data set.


4.3 Model 3

For Model 3, our goal was to add any other variables in the data set that we have sufficient detail about to interpret and that are not strongly covariant with any of the variables already in Model 2. It is not our expectation that adding these variables will substantially improve the predictive value of the model, but it will help us to evaluate how close Model 2, which explains nearly 81% of the variance in crime rate, is to the boundary of our data set's explanatory power.

For this model, we additionally consider whether or not a county is in the central part of the state (**central**), the ratio of prison sentences to convictions (**prbpris**), and the average prison sentence following conviction for a crime (**avgsen**). Given that counties in the center of the state include some of the largest cities (e.g., Raleigh, Durham, Greensboro, Charlotte) and cities tend to have more crime than rural areas, we expect **central** to have a positive relationship with **crmrte**. As in Model 2, we expect the variables related to punishment for crimes, such as **prbpris** and **avgsen**, to also be negatively associated with **crmrte**; as punishments increase in their likelihood and severity, people should be less likely to engage in criminal activity.

One variable we do not include in this model is **urban**, or whether a county is designated as predominantly urban by the U.S. Census. This is because our preliminary analysis suggested **urban** is highly correlated with **density**, which is one of our key explanatory variables. The scatterplot below shows the strength of this relationship:

```
plot(df_clean$urban , df_clean$density,
     main= "Population Density vs Urban Status",
     ylab= "Population Density",
     xlab= "Urban Status", cex.main = 1, cex.lab=.75, cex.axis = .75)
abline(lm(density ~ urban, data = df_clean))
```




We build the model and evaluate its consistency with the remaining statistical assumptions. The detailed results can be seen in **Table 3**. We observe that none of the new variables have statistically significant coefficients, nor does the predictive value of the model, as measured by the adjusted R^2 , change relative to Model 2.

```
# Build Model 3
model_3 = lm(crmrte ~ density + taxpc + pctymle
             + west + polpc + prbarr + prbconv + pctmin80
             + central + avgsen + prbpris,
             data = df_clean)
```

We briefly examine the model with regard to the remaining assumptions.

```
# Evaluate influence in Model 3
plot(model_3, which= 5,
     cex.main = 1, cex.lab=.75, cex.axis = .75)
```




(crmrate ~ density + taxpc + pctymle + west + polpc + prbarr + prbcor

We note that point 25, representing Dare county, is still exhibiting a Cook's distance of greater than 1.

We test Assumption 3 by evaluating whether any of the variables display perfect collinearity with any of the others, and find that none do.

To validate Assumption 4, exogeneity, we check the plot of residuals against fitted values and observe a fairly flat mean across the entire range of fitted values.

```
plot(model_3, which = 1,
      cex.main = 1, cex.lab=.75, cex.axis = .75)
```



(crmte ~ density + taxpc + pctymle + west + polpc + prbarr + prbcor

Assumptions 5 and 6 are validated for this model just as they were for Models 1 and 2.

Statistical Inference and Interpretation for Model 3

```
coeftest(model_3, vcov = vcovHC)
```

```
##
## t test of coefficients:
##
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) 2.4830e-02 9.8357e-03  2.5245 0.013617 *
## density     5.8503e-03 1.4385e-03  4.0670 0.000113 ***
## taxpc       1.4705e-04 2.4494e-04  0.6004 0.550007
## pctymle     7.5521e-04 3.6326e-04  2.0790 0.040909 *
## west       -4.9285e-03 3.4884e-03 -1.4128 0.161693
## polpc       6.8611e-03 2.0912e-03  3.2809 0.001548 **
## prbarr     -5.4372e-04 1.3004e-04 -4.1812 7.520e-05 ***
## prbconv    -1.8259e-02 4.1137e-03 -4.4385 2.942e-05 ***
## pctmin80    2.6353e-04 8.9263e-05  2.9523 0.004166 **
## central    -3.9980e-03 2.6039e-03 -1.5354 0.128735
## avgsen     -2.9651e-04 3.7968e-04 -0.7809 0.437200
## prbpris     3.7196e-05 1.3020e-04  0.2857 0.775875
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


Our heteroskedasticity-robust t -test shows that none of the added variables are statistically significant at a .05 level. This means we cannot reject the null hypotheses for **central**, **avgsen**, and **prbpris** as they relate to **crmrte**. In practical terms, this means that whether a county is in the central part of the state does not effectively predict that county's crime rate, nor does it seem to matter what the likelihood or severity of prison sentences are. Given this evidence, we opt not to make policy recommendations based on these variables.

4.4 Model 4

For this model, we include every variable in the data set. This is done to set an upper limit on the possible R^2 , or predictive power of the data at hand. The resulting model is not a parsimonious one, and thus will not be used as the basis for any policy recommendations.


```
# Build Model 4
model_4 = lm(crmrte ~ density + taxpc + pctymle
             + west + polpc + prbarr + prbconv + pctmin80
             + central + avgsen + prbpris
             + urban + wcon + wtuc + wtrd + wfir + wser
             + wmfg + wfed + wsta + wloc + mix,
             data = df_clean)
```

We refer the reader to **Table 3** for the details of this model, but note that the adjusted R^2 remains at 81%, the same as Model 2. In the leverage plot below, the points for Dare county and Warren county display Cook's values of 1 or more. Otherwise, we will not evaluate this model.

```
# Evaluate influence plot
plot(model_4, which = 5,
     cex.main = 1, cex.lab = .75, cex.axis = .75)
```

```
## Warning in sqrt(crit * p * (1 - hh)/hh): NaNs produced
```

```
## Warning in sqrt(crit * p * (1 - hh)/hh): NaNs produced
```



```
(crmrte ~ density + taxpc + pctymle + west + polpc + prbarr + prbconv)
```

Statistical Inference and Interpretation for Model 4

```
coeftest(model_4, vcov = vcovHC)
```

```
##
## t test of coefficients:
##
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) 1.3853e-02 3.0755e-02  0.4504 0.6538622
## density      5.3314e-03 1.4895e-03  3.5793 0.0006464 ***
## taxpc        1.6240e-04 2.8408e-04  0.5717 0.5694537
## pctymle      1.0125e-03 4.7826e-04  2.1170 0.0379748 *
## west        -2.5652e-03 4.4698e-03 -0.5739 0.5679579
## polpc        6.9679e-03 2.9536e-03  2.3591 0.0212406 *
## prbarr       -5.1466e-04 1.5689e-04 -3.2805 0.0016467 **
## prbconv      -1.8633e-02 6.5853e-03 -2.8295 0.0061464 **
## pctmin80     3.2542e-04 1.3849e-04  2.3497 0.0217429 *
## central     -4.2416e-03 3.7423e-03 -1.1334 0.2610725
## avgscen     -3.9858e-04 5.5361e-04 -0.7200 0.4740570
## prbpris      3.1727e-05 1.3586e-04  0.2335 0.8160642
## urban       -9.6498e-05 8.2752e-03 -0.0117 0.9907307
## wcon         2.3025e-05 3.2876e-05  0.7004 0.4861334
## wtuc         6.1914e-06 1.9862e-05  0.3117 0.7562178
## wtrd         2.8767e-05 8.7294e-05  0.3295 0.7427756
## wfir        -3.5455e-05 3.5699e-05 -0.9932 0.3242068
## wser        -1.7158e-06 9.9447e-05 -0.0173 0.9862856
## wmfgr        -8.9675e-06 1.7469e-05 -0.5133 0.6094087
## wfed         2.9075e-05 3.7780e-05  0.7696 0.4442480
## wsta        -2.2302e-05 3.6828e-05 -0.6056 0.5468431
## wloc         1.4456e-05 8.5367e-05  0.1693 0.8660410
## mix         -1.8693e-04 2.2922e-04 -0.8155 0.4176761
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

As with Model 3, we see that none of the new variables are statistically significant at the .05 level. More importantly, we observe that the statistical significance of several of our Model 1 and 2 variables – including **polpc**, **prbarr**, **prbconv**, **pctmin80**, **pctymle**, and **density** – has decreased as a result of these new variables being added. This suggests that the effect of these established variables has been absorbed by the new variables, although the mechanism(s) of this change are not transparent.

Before concluding our statistical analysis for Model 4, we think it would be useful to explore whether any of the explanatory variables are jointly significant with regard to the crime rate. This can be determined using an *F*-test. Specifically, we will create a restricted model containing only the statistically significant covariates, and compare this restricted model to the unrestricted Model 4 to see if the combined effect of Model 4's individually insignificant variables changes to become statistically significant.

```
restricted_model <- lm(crmrte ~ density + pctymle + polpc
+ prbarr + prbconv + pctmin80, data = df_clean)
waldtest(model_4, restricted_model, vcov=vcovHC)
```

```
## Wald test
##
## Model 1: crmrte ~ density + taxpc + pctymle + west + polpc + prbarr +
##          prbconv + pctmin80 + central + avgscen + prbpris + urban +
##          wcon + wtuc + wtrd + wfir + wser + wmfgr + wfed + wsta + wloc +
##          mix
## Model 2: crmrte ~ density + pctymle + polpc + prbarr + prbconv + pctmin80
```

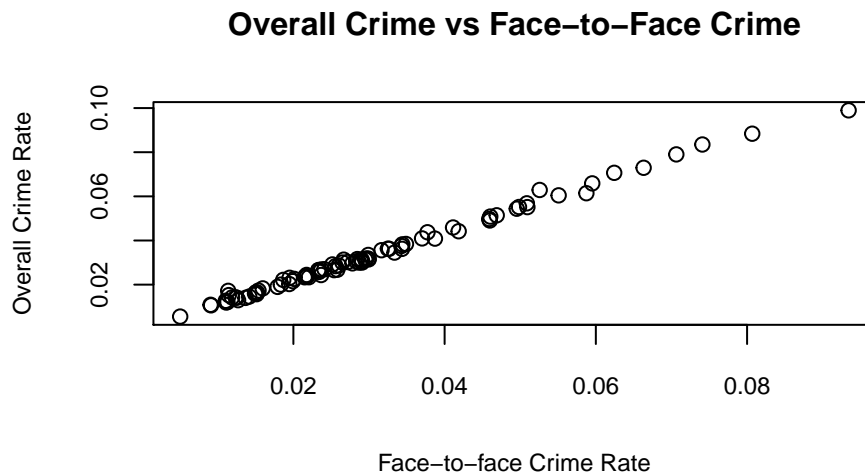
```
##   Res.Df  Df       F Pr(>F)
## 1      67
## 2      83 -16 0.8621 0.6131
```

The F -test reveals that even the combined statistical significance of the 16 additional variables in the unrestricted Model 4 ($p = 0.61$) is not enough to suggest their joint statistical significance. As such, we fail to reject the null hypothesis that these additional factors have no association with crime rate.

4.5 Model 5

Lastly, we return to our face-to-face crime rate and attempt a model using it as the outcome variable. As previously noted, the purpose of this model is to determine the extent to which Model 1, intended to predict the general crime rate, effectively extends to predict the face-to-face crime rate. Our expectation is the the results for Model 5 and Model 1 will be very similar, as `crmrate` and `ftfcrmrate` are closely associated:

```
plot(crmrate ~ ftfcrmrate, data = df_clean,
     main = "Overall Crime vs Face-to-Face Crime",
     xlab = "Face-to-face Crime Rate",
     ylab = "Overall Crime Rate",
     cex.main = 1, cex.lab = .75, cex.axis = .75)
```




```
# Build Model 5
model_5 <- lm(ftfcrmrate ~ density + taxpc + pctymle, data=df_clean)
```

Implementation of Model 5 shows that population density, tax per capita, and the percentage of young males in the population explain approximately 62% of the variation in the face-to-face crime rate; this is almost identical to their predictive power relative to general crime rate. The results of this analysis are summarized in **Table 3**.

Below, we see that the leverage plot for Model 5 is virtually identical to that for Model 1.

```
# Evaluate influence plot
suppressWarnings(plot(model_5, which = 5))
```



Statistical Inference and Interpretation of Model 5

```
coefTest(model_5, vcov=vcovHC)
```

```
##
## t test of coefficients:
##
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.00731872  0.01078922 -0.6783  0.499380
## density      0.00698061  0.00107721  6.4803 5.497e-09 ***
## taxpc        0.00035398  0.00026346  1.3436  0.182615
## pctymle      0.00169822  0.00062913  2.6993  0.008364 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The t -test on coefficients for Model 5 closely resembles that of Model 1, with **density** and **pctymle** showing statistical significance at the .05 level, and **taxpc** failing to demonstrate statistical significance. We note with interest that **pctymle** has greater statistical significance in this model than it did in Model 1, suggesting that the percent of young males in a county has a stronger relationship with the rate of personal crime than for the rate of crime in general.

Given that we see very similar practical and statistical significance for the covariates in Models 1 and 5, the answer to our corollary question, *What are the variables associated with the **face-to-face** crime rates across counties in North Carolina?*, appears to be “The same variables that are associated with the general crime rate, and the impact is generally on a comparable scale.” That being said, it does appear that the percentage of young males may have a special relationship with the rate of face-to-face crime that would be a useful topic of future research.

As noted previously, our experience indicates that constituents are interested in policies and programs intended to reduce crime, but are especially concerned with those directed at reducing personal crime. Model 5 demonstrates to the candidate that policies and programs intended to reduce the general crime rate, about which we have made a number of recommendations, are also likely to reduce the personal crime rate. What matters most in this case is constituents’ *perception*: The candidate is likely to get more support for policies and programs that are framed in terms of reducing personal crime, particularly violent crime. Moreover, the candidate can promote policies and programs in terms of reducing personal crime while knowing that implementation of these policies should also reduce the general crime rate. This analysis provides a foundation

for the candidate to pursue this strategy without sacrificing veracity or efficiency.

4.6 Model Summaries

Table 3 summarizes the results of all five of the presented models. Model 2 appears to be the best balance of predictive ability and parsimony, given that it has the highest adjusted R^2 and the lowest AIC value.

```
model_1$AIC <- AIC(model_1)
model_2$AIC <- AIC(model_2)
model_3$AIC <- AIC(model_3)
model_4$AIC <- AIC(model_4)
model_5$AIC <- AIC(model_5)

stargazer(model_1, model_2, model_3, model_4, model_5, type = "latex",
  se = list(sqrt(diag(vcovHC(model_1))), sqrt(diag(vcovHC(model_2))),
    sqrt(diag(vcovHC(model_3))), sqrt(diag(vcovHC(model_4))),
    sqrt(diag(vcovHC(model_5)))),
  title = "Linear Models Predicting Crime Rate",
  keep.stat = c("aic", "rsq", "n", "adj.rsq"),
  no.space = TRUE, table.placement = "!", header = FALSE,
  star.cutoffs = c(0.05, 0.01, 0.001))
```

4.7 Reservations & Caveats

Although we are generally pleased with the outcome of our analysis, we note several reservations we have about the data and the findings, along with caveats we believe the reader should consider when interpreting the results and evaluating our recommendations.

First, we remind the reader that the data we were provided may not comprise a simple random sample of any broader population. We do not know if the counties included in the data set were truly randomly selected from among all counties in North Carolina in 1987, nor can we be confident that the measurements taken for the respective counties are independent. We can imagine, for instance, that geographically close counties, could have mutually dependent crime rates.

Second, the most commonly used statistical methods are designed for samples that comprise less than 10% of the population, whereas our data set includes 90% of the population of North Carolina counties in 1987. Given that we have nearly sampled the entire population, the usual business of statistical inference – trying to learn about a population based on a sample – is not completely necessary in order to make meaningful claims about the relationships among crime rate and other variables in the current case. Even supposing the relationships we discovered *only* hold for the 90 counties that we have data for, both the candidate and their constituents are likely to be satisfied with policies and programs that effectively reduce crime in these counties alone.

Third, the context of standard errors and hypothesis tests is a hypothetical scenario in which we imagine our data as a random sample of counties drawn from some theoretical infinite population of possible sets of measurements of each variable. Rather than being merely an abstract exercise in multiverse theory, this approach allows us to consider how noise in our data might or might not be obscuring some “true” relationship between the explanatory factors we considered and the outcome of crime rate.

Finally, as our models were created after an initial exploration of the data, the evidence of our hypothesis tests is best viewed as a jumping-off point for future research. Were we given new data from another year or another state, or even data that were drawn randomly from multiple years or states, we could test hypotheses based on our current findings (e.g., that change in population density has a nonzero effect on crime rate). Even better would be the possibility to implement policy changes in some North Carolina counties but not

Table 3: Linear Models Predicting Crime Rate

	<i>Dependent variable:</i>				
	crmte				ftfcrmte
	(1)	(2)	(3)	(4)	(5)
density	0.008*** (0.001)	0.005*** (0.001)	0.006*** (0.001)	0.005*** (0.001)	0.007*** (0.001)
taxpc	0.0004 (0.0003)	0.0002 (0.0002)	0.0001 (0.0002)	0.0002 (0.0003)	0.0004 (0.0003)
pctymle	0.002* (0.001)	0.001* (0.0004)	0.001* (0.0004)	0.001* (0.0005)	0.002** (0.001)
west		-0.002 (0.003)	-0.005 (0.003)	-0.003 (0.004)	
polpc		0.007** (0.002)	0.007** (0.002)	0.007* (0.003)	
prbarr		-0.001*** (0.0001)	-0.001*** (0.0001)	-0.001** (0.0002)	
prbconv		-0.019*** (0.004)	-0.018*** (0.004)	-0.019** (0.007)	
pctmin80		0.0003*** (0.0001)	0.0003** (0.0001)	0.0003* (0.0001)	
central			-0.004 (0.003)	-0.004 (0.004)	
avgsen			-0.0003 (0.0004)	-0.0004 (0.001)	
prbpris			0.00004 (0.0001)	0.00003 (0.0001)	
urban				-0.0001 (0.008)	
wcon				0.00002 (0.00003)	
wtuc				0.00001 (0.00002)	
wtrd				0.00003 (0.0001)	
wfir				-0.00004 (0.00004)	
wser				-0.00000 (0.0001)	
wmfg				-0.00001 (0.00002)	
wfed				0.00003 (0.00004)	
wsta				-0.00002 (0.00004)	
wloc				0.00001 (0.0001)	
mix				-0.0002 (0.0002)	
Constant	-0.009 (0.013)	0.019* (0.008)	0.025* (0.010)	0.014 (0.031)	-0.007 (0.011)
Observations	90	90	90	90	90
R ²	0.640	0.824	0.830	0.855	0.631
Adjusted R ²	0.628	0.807	0.806	0.807	0.619
Akaike Inf. Crit.	-542.122	-596.430	-593.615	-585.586	-556.107

Note:

*p<0.05; **p<0.01; ***p<0.001

others, creating an experiment of sorts. However, we advise that the general public might not take kindly to this latter approach.

5. Omitted Variables

Although our Models 1 and 2 do provide some useful information that can inform government policy, there are a number of omitted variables that we did not have access to in this analysis that we suspect have meaningful associations with the crime rate. We must attempt to identify these variables and deduce the impact we believe they would have, or else we risk biasing our conclusions by considering only the variables we can measure.

We expect that **household earnings** would be negatively correlated with crime, as wealthier areas typically have less crime. We would expect correlation between household earnings and density to be positive, because salaries are typically higher in cities. Therefore, omitted variable bias is negative, and if we controlled for household earnings, the fitted values for a given population density value would likely be higher. The negative omitted variable bias on a factor with a positive coefficient would make the predictive power of the affected variable greater.

Relatedly, we believe that **socioeconomic diversity** is likely to have a strong association with crime rate. If a county has a mix of people who have ample resources and those who have very little, the disparity can create social tension. Furthermore, if those who do not have resources live in close proximity to those who have many, those without may be motivated to engage in crime to secure the resources they need from those who have them. We could measure socioeconomic diversity by measuring the gap between the 1st and 3rd quartiles of household earnings (see below). We would expect a large gap to be associated with a high crime rate, and we would also expect a positive correlation between our measured income gap and population density, as dense urban areas tend to have both wealthy and impoverished people living in close proximity. In this case, omitted variable bias is positive, and the fitted values would be lower for a given population density value if we had socioeconomic diversity as a variable. This would diminish the effect of population density, bringing the coefficient closer to zero.

Connected with both earnings level and socioeconomic status is the **unemployment rate**, which we think would likely have a positive association with the crime rate. When people are unable to find or keep a job, they are likely to struggle to cover their living expenses or they may find themselves seeking meaningful ways to spend their time. We can imagine both of these scenarios contributing to an increase in the crime rate. We expect the correlation between the unemployment rate and percent of young males to be positive, as many young men do not participate in the labor force. Therefore, the omitted variable bias is positive, and the fitted values would be lower for a given percent young male value if we had data on the unemployment rate. This would lower the effect of percent young male on crime rate, bringing the coefficient closer to zero.

People's ability to secure successful employment is likely to be associated with **educational attainment**, and thus this variable also seems likely have potential impact on the crime rate. If we added highest education level to a model, we would expect its coefficient with regard to crime to be negative; when people have more education, they are more likely to have an adequate income and possess skills or abilities that facilitate their contributing meaningfully to society. We anticipate a positive correlation between education level and population density, as urban areas tend to have higher education levels due to job opportunities and the presence institutions of higher education. Therefore, the omitted variable bias for education level is negative, and the fitted values for a given value of population density would likely be higher if we could control for education level. That is, a negative omitted variable bias on a positive coefficient would only increase the coefficient, making the effect of population density even greater.

Beyond factors related to people's ability to provide for their living needs without engaging in crime, we think there are several important community-centered variables to be considered. First, **infrastructure quality** would be negatively correlated with crime, as areas with well developed infrastructure make it easier for people to live, work, and play with a minimum of environmental stress. We would expect a positive correlation between infrastructure quality and tax per capita as taxes are used to fund the infrastructure development. Therefore, the omitted variable bias for infrastructure quality is negative, and the fitted values

for a given value of tax per capita would be higher if we controlled for infrastructure quality. Second, we expect **community membership** to be negatively correlated with the crime rate, as a sense of belonging in one's community seems likely to reduce the desire to harm that community through negative behaviors like engaging in crime. We could expect a negative correlation between community membership and the percent of young males, given that young males are typically moving through a city and not planting roots. Perhaps the more obvious correlation is a negative one with population density, because community membership is harder to develop when there are many people to know. Either way, the omitted variable bias for community membership would be positive. We expect that the fitted values for a given value of percent young male or population density could be lower if we controlled for community membership, making the effect of percent young male less extreme, closer to zero. We note also that community membership would be a challenging variable to measure accurately. Third, we believe **negative police interaction** to be positively correlated with crime, and positively correlated with police per capita, leading to a positive omitted variable bias. This leads to fitted values for a given value of police per capita to be lower, if we controlled for negative police interaction, making the effect of police per capita less extreme.

Somewhat connected to both the wealth of an area and the strength of community in an area is the omitted variable of **tourism**. We expect tourism to be negatively correlated with crime rate, given that it tends to strengthen the local economy, leading to less crime. We also expect tourism to be positively correlated with **west** and the coastal region (including counties like Dare), as these areas have lots of tourist activity in the way of national parks and other attractions. This would lead to a negative omitted variable bias. If we controlled for tourism, we expect the fitted values for a given value of west to be higher, making the effect of **west** closer to zero.

Lastly, we believe cultural values like **religious faith** may play a role in the crime rate. For example, if the people living in a particular region report a strong religious faith or high levels of religious activity, they may be less likely to engage in crime, because many faiths have specific prohibitions against engaging in violence, stealing, damaging what belongs to others, and other illegal behavior. In other words, higher levels of religious beliefs or activity should negatively correlated with crime rate. Moreover, we suspect that the more uniform the religious faith in a region, the more robust this effect should be. This is because religious differences in a community can create social conflict, and such conflict may motivate crime against those perceived as "other". We note that shared cultural values, such as religious faith, may represent a form of **community membership**, discussed above.

We note that this discussion of omitted variables is speculative and should be reinforced with research. Ideally, this research would include randomized, controlled trials where possible.

6. Conclusion

Based on our regression analysis of the available data, we find that population density, the percentage of young males, police per capita, probability of arrest, probability of conviction, and percent minority population explain the majority of the variation in crime rate in North Carolina in the late 1980s. However, we suspect several endogenous variables' effects to be contained in those variables, such as community membership, unemployment rate, and educational attainment.

Based on this analysis, we have several policy suggestions applicable to candidates seeking or defending office in North Carolina. Specifically, candidates are advised to pursue programs that increase public access to key resources, encourage community involvement, and decrease perceptions of social injustice. Moreover, we encourage the candidates to discuss their crime reduction plans from the perspective of reducing personal crime, including violent crime, because such crimes are typically of greater concern to constituents than property crimes and are likely to generate voting behavior that favors the candidate.

Although we believe our analysis supports our recommendations well, it does have some important limitations. First, the data are reported by county, which is a rather coarse measure for examining crime rate and the variables that influence it. It seems unlikely, for example, that population density and tax per capita are evenly distributed across a single county, so potentially meaningful variations in these variables are not transparent at the county level. Similarly, some major metropolitan areas stretch across multiple counties,

and their police forces and court systems are accustomed to interacting with and accommodating one another. As such, the behaviors of the police and courts might appear more similar in these counties, when this is actually an artifact of the city size. Second, our data set is limited to North Carolina counties in 1987. Not only is this data more than 30 years old, but it is restricted to a specific state at a very specific time. As such, we cannot generalize our findings to other states or years, or even test any hypotheses we might generate. Third, we have limited information about how the data were collected, recorded, or prepared prior to being delivered to us. In the absence of this information, we cannot assess the quality of the data or develop clear strategies for addressing in weaknesses that may have resulted from the collection methodology. Lastly, the counties that are omitted from the data set appear to be those with lower population. Given that population density has proven to be an important variable in our models, it would be useful to have data from these counties to make sure our understanding of this variable is as complete and accurate as possible.